

The 5 Stages of Scale

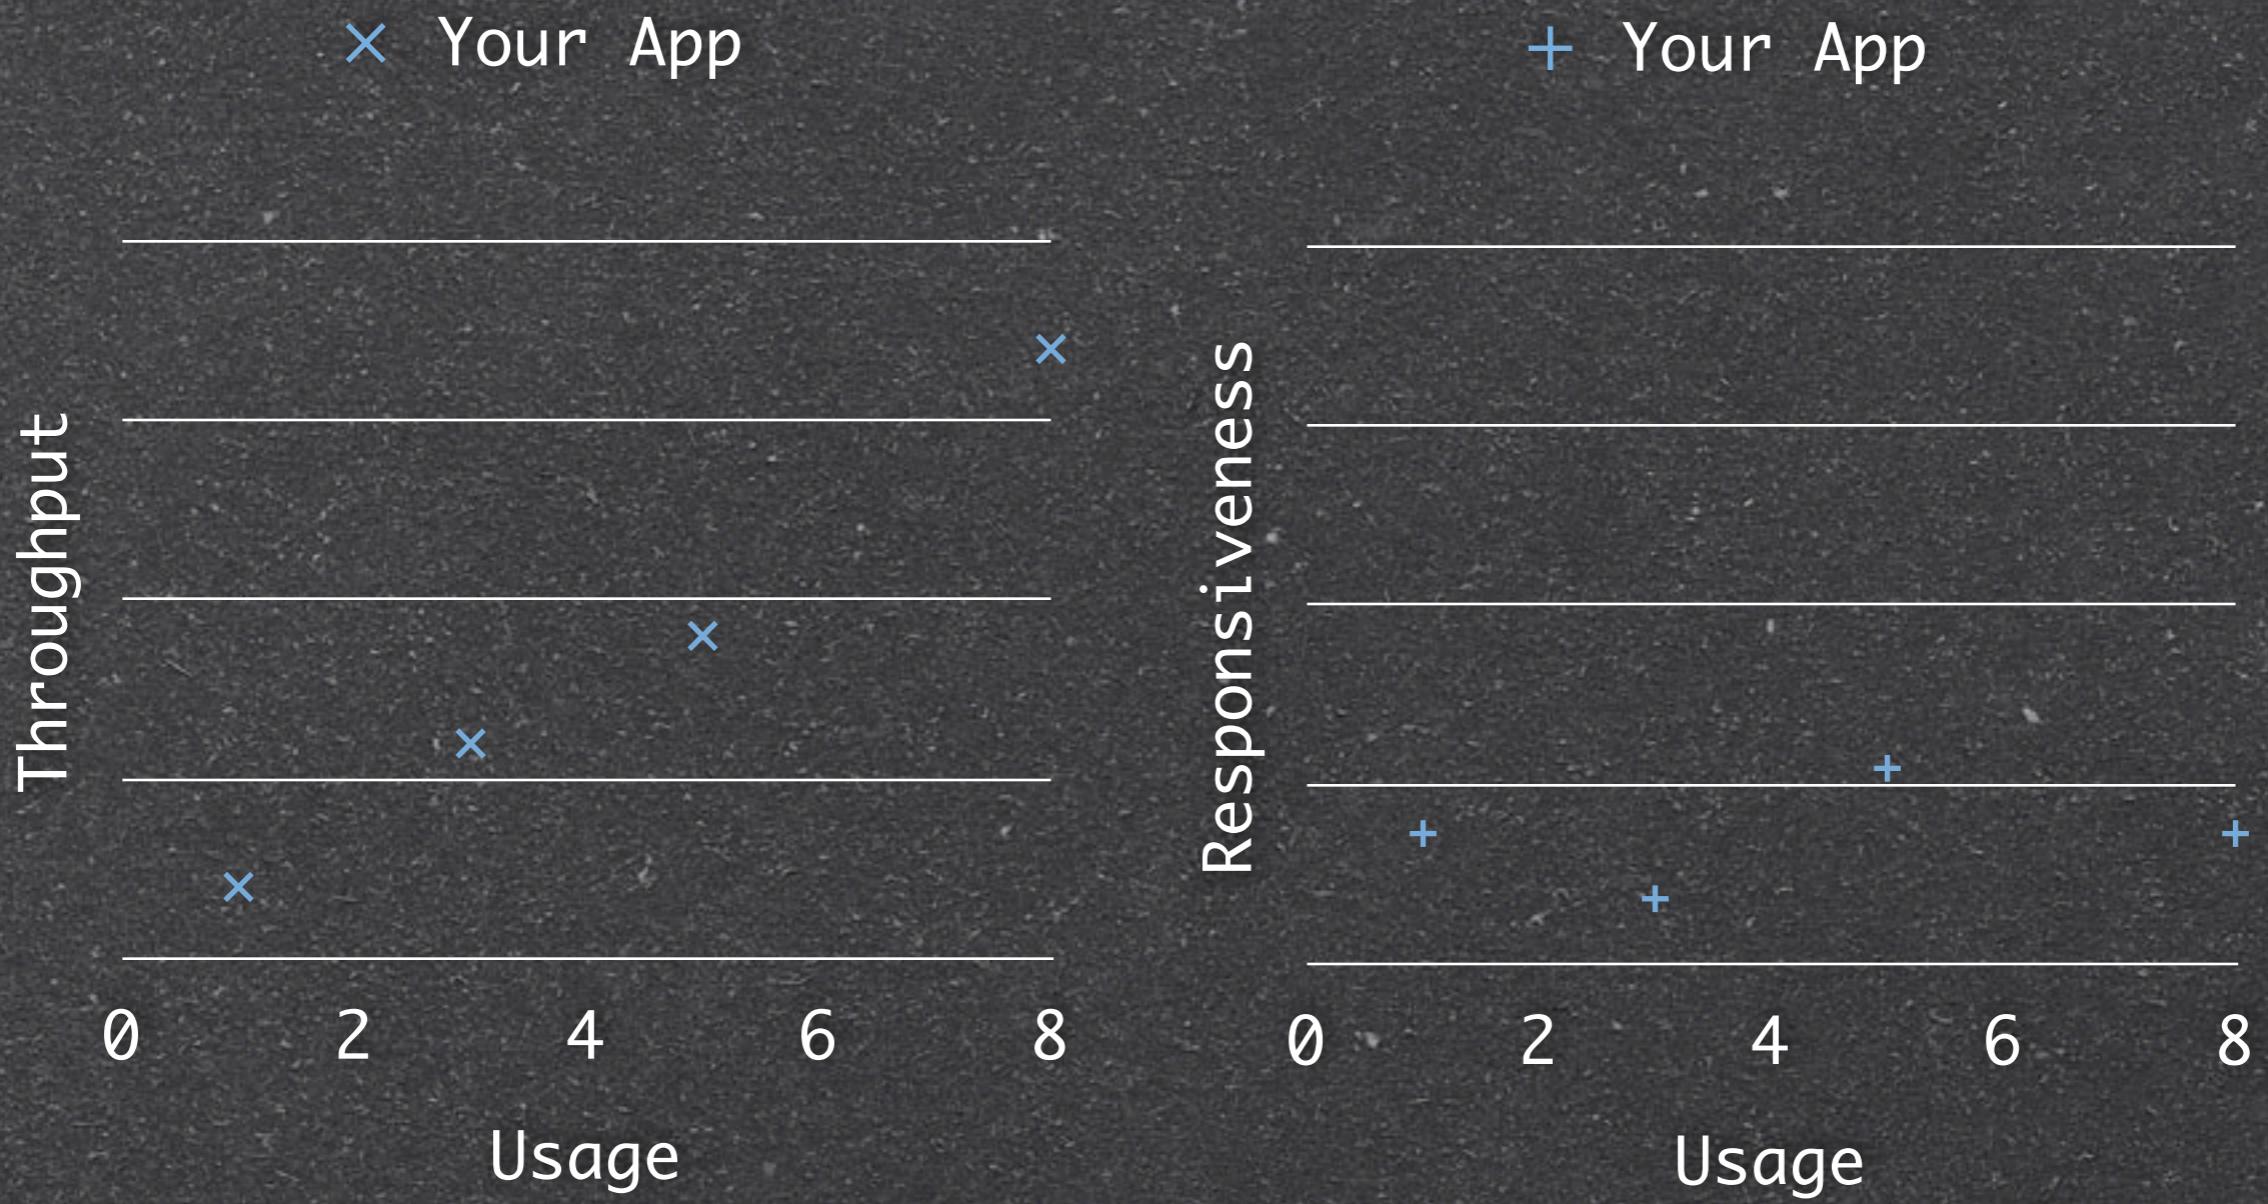
Christopher Smith

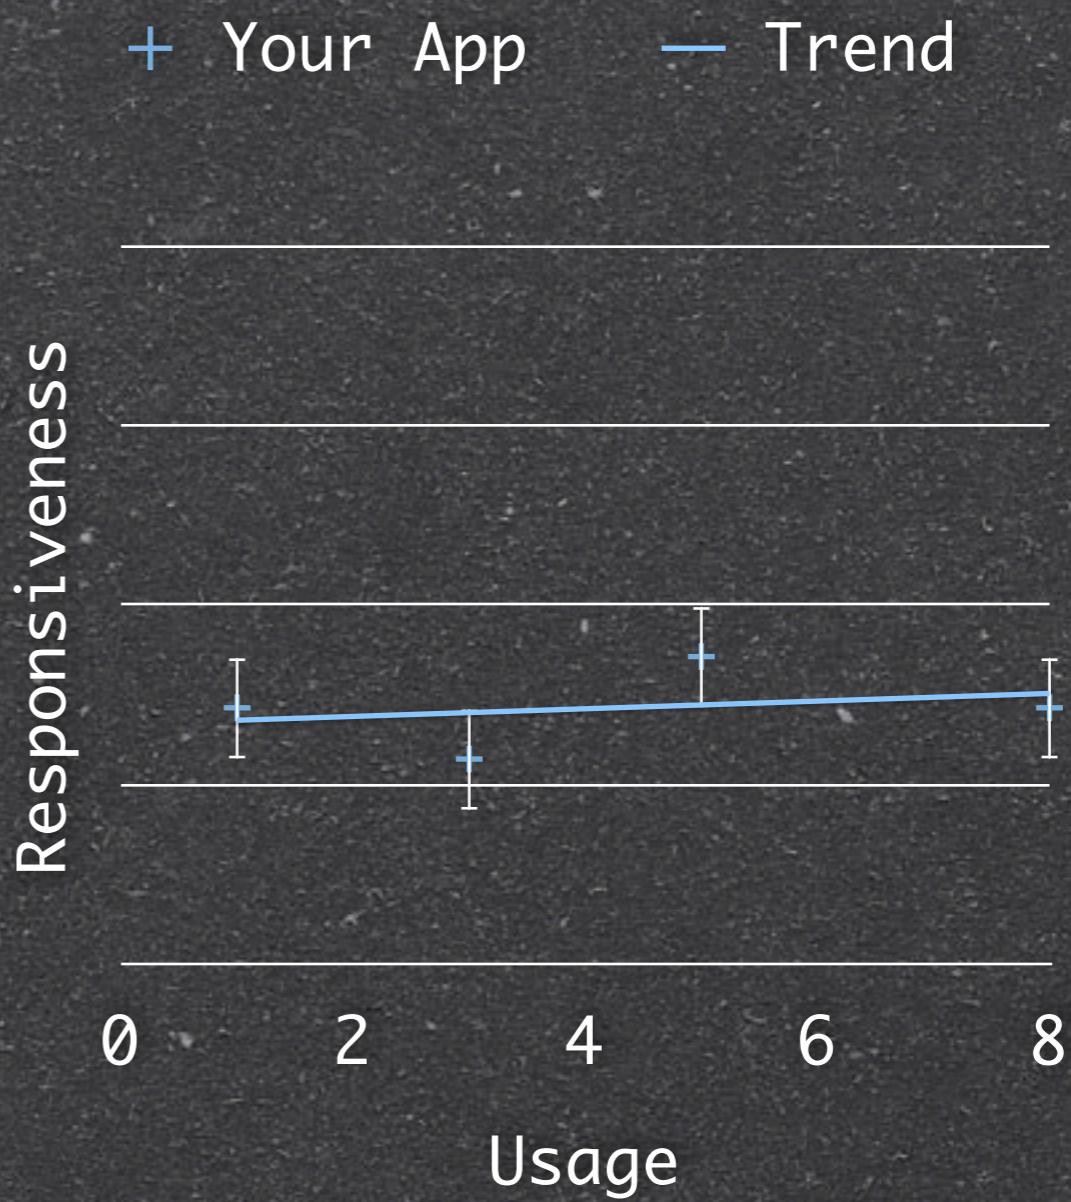
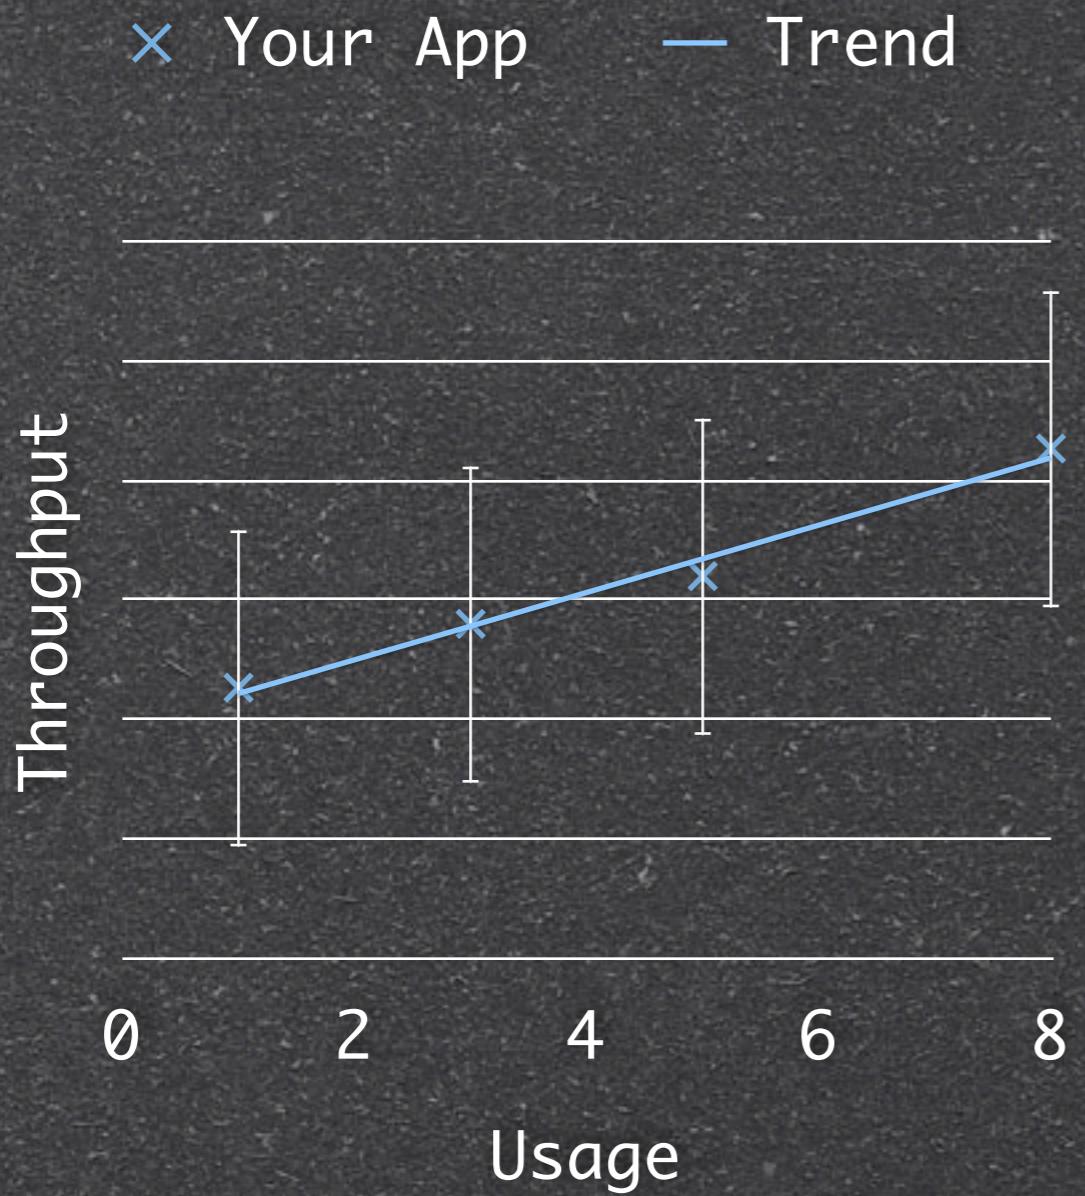
Who am I?

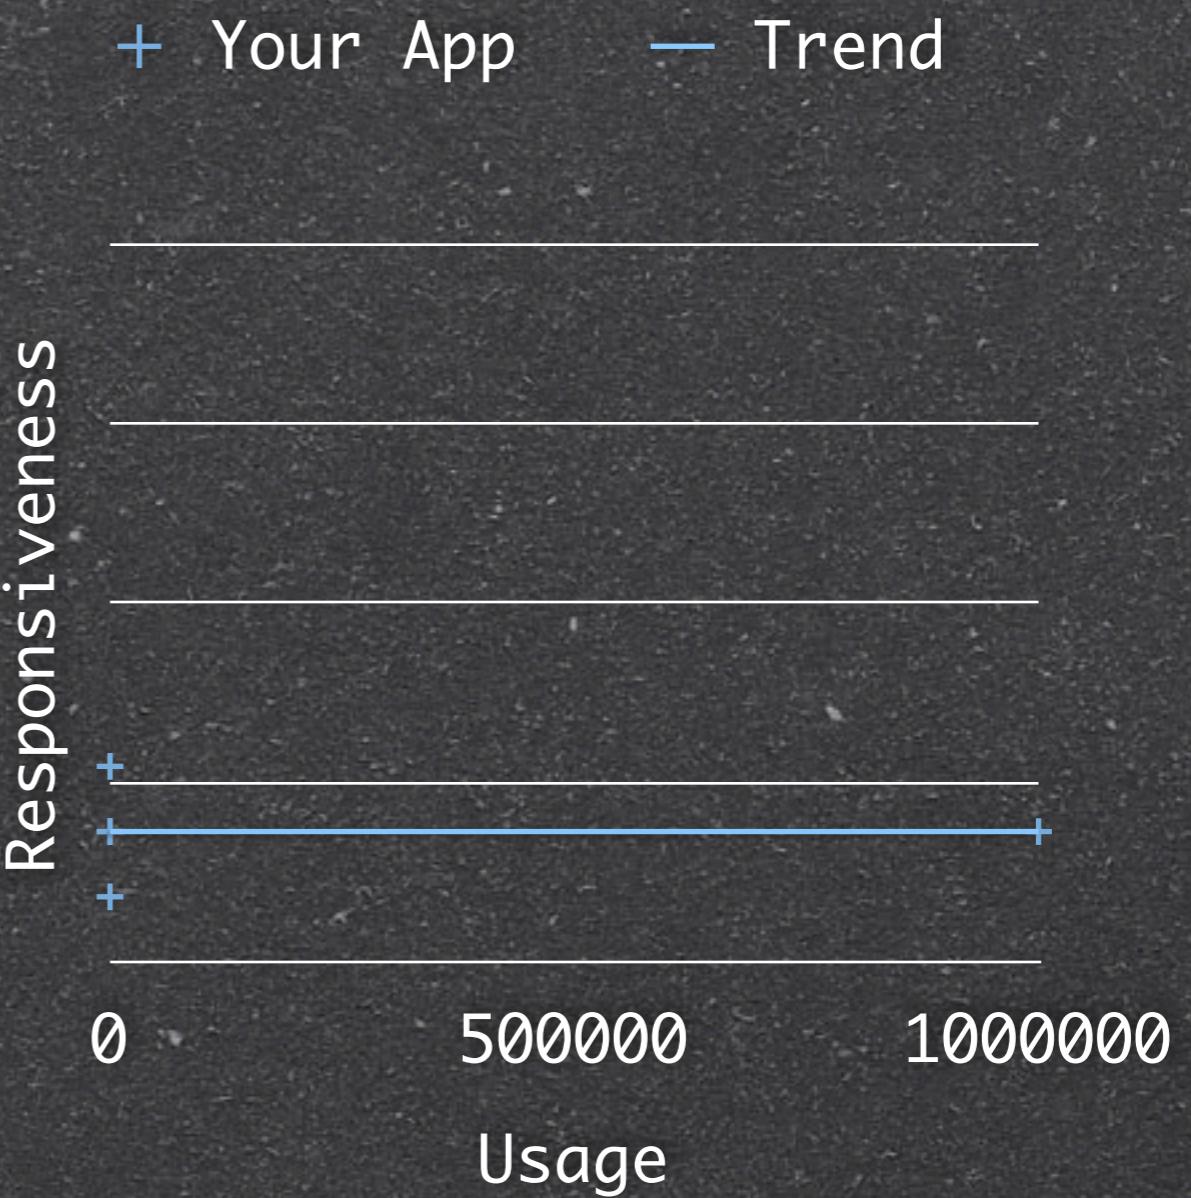
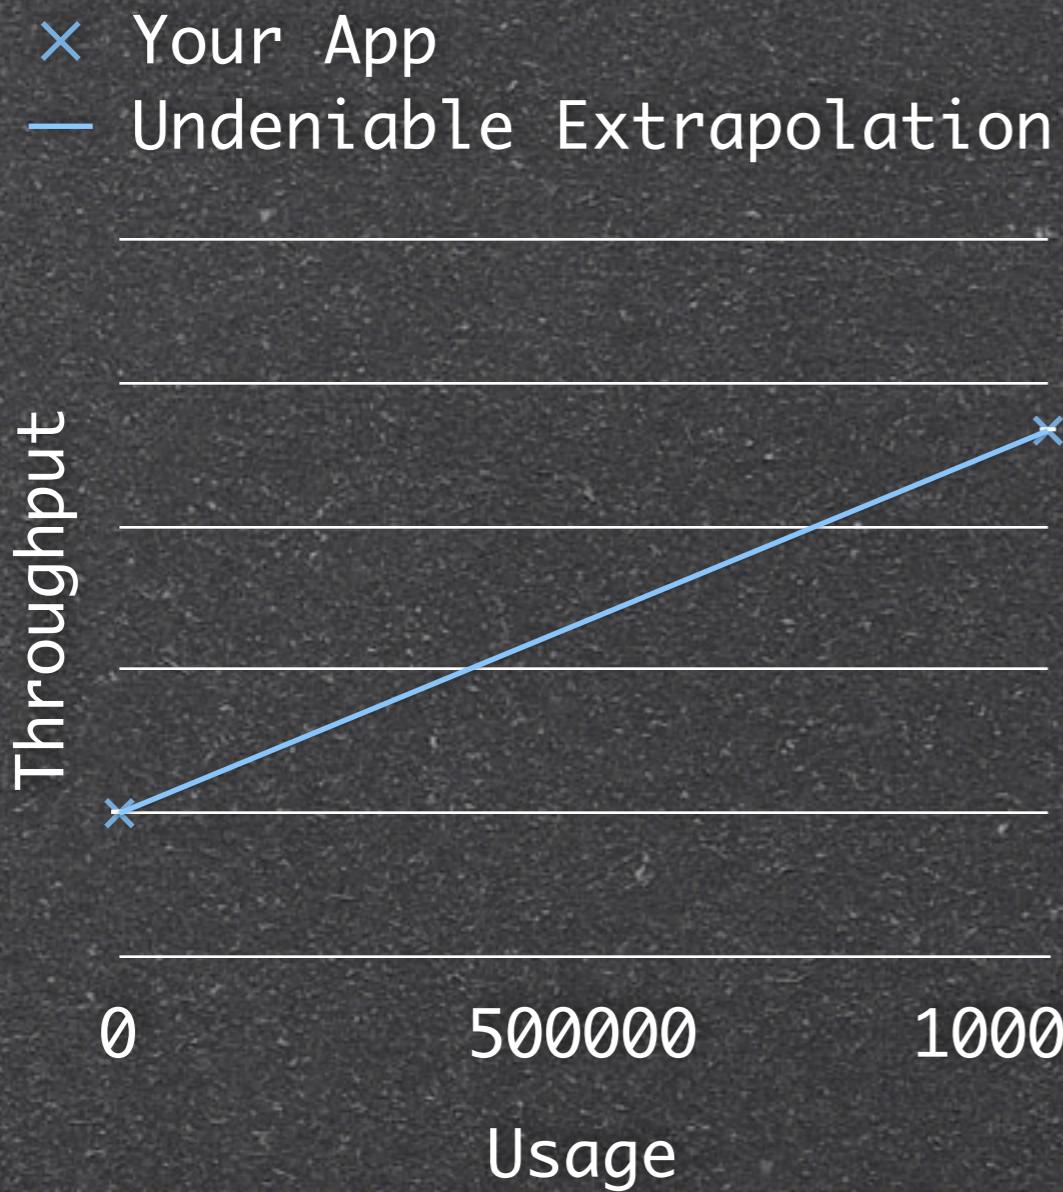
- Two decades experience
- Half of that in online advertising
- Internet systems engineering
- Scaling web serving, data collection & analysis
- Places big & small.

Scalability

- scale: *v.tr.*
 - 1. To clear or strip of scale or scales.
 - 2. Weigh a specified weight.
 - 3. Climb up or over (something steep)







✖ Your App
— Undeniable Extrapolation

+ Your App — Trend



Scalability Envelopes

- There is always a “next” bottleneck.
- In case of scalability problem...
- 6 envelopes

Envelope 0

- Session partitioning
- Commodity: load balancer, multi-*
- Linear scale for CPU
- Limit: C10K?

Envelope 1

- Read Caching
 - Reverse-proxy
 - memcached
 - CDN
- $\log(n)$ scale: thank you Zipf
- Limit: ~200 w/sec

Envelope 2

- Get a real persistence framework
 - Data structures FTW!
 - DB: concurrent read/write
 - MOM: queuing/event I/O/TP monitors
 - Cheat on ACID (particularly C & D)
- log(n) scale?
- 1000-10000 w/sec

Tipping over



Scaling Catamaran's

- RAM caching I/O
- RAID
- Threads (sometimes)
- Packet loss (UR DURING IT WRONG)
- SSD's?

Jeff Dean's Numbers

- Latency Comparison Numbers

• -----

• L1 cache reference	0.5	ns		
• Branch mispredict	5	ns		
• L2 cache reference	7	ns	14x L1 cache	
• Mutex lock/unlock	25	ns		
• Main memory reference	100	ns	20x L2 cache, 200x L1 cache	
• Compress 1K bytes with Zippy	3,000	ns		
• Send 1K bytes over 1 Gbps network	10,000	ns	0.01 ms	
• Read 4K randomly from SSD*	150,000	ns	0.15 ms	
• Read 1 MB sequentially from memory	250,000	ns	0.25 ms	
• Round trip within same datacenter	500,000	ns	0.5 ms	
• Read 1 MB sequentially from SSD*	1,000,000	ns	1 ms	4X memory
• Disk seek	10,000,000	ns	10 ms	20x datacenter roundtrip

Problem: IO latency

- Throughput: 2x every 18 months
- Latency:
 - CPU: <2x every 18 months
 - LAN network: 2x every 2-3 years
 - Memory: 2x every 3-5 years
 - Disk: 2x every decade? (SSD?)
 - WAN Network: 1x every...

Problem IO Latency

- Traditional indexes on the wrong side
 - Turns a scan in to a seek
 - Index lookup: scan 0.1% of records + 1 random seek
 - Scan: scan 100% of records, 0 random seek
 - Seek is 10ms & Scan is 100Hz -> 10x win
 - Seek is 1ms & Scan is 1GHz -> 1000x loss

Envelope 3

- Real partitioning of IO
- Move code, not data
- Commodities: Map/Reduce (Hadoop), DHT (Cassandra, HBase, Riak)
- CAP Theory limiting sync'ing

Envelope 4

- Route new data through data partitions
 - Using MOM/EventIO “the right way”
 - ESP/CEP: Eigen, Storm, Esper, StreamBase, Ømq, etc.

Envelope 5

- Cheat more on reliability.
 - UDP w/o reliability > TCP
 - Measure loss vs. prevent loss
 - Horseshoes, hand grenades, features...?

Integrated Systems

- Combined IO management solutions:
 - real-time memory key/value lookup
 - LSM + bitmap indexes + etc.
 - eventual consistency
 - mobile code for batch processing
- Cassandra, HBase, etc.

Efficient Logging

- Events in efficient machine parseable form: (protobuf, thrift, etc.)
- Event source writes only to NIC
- UDP Multicast
- Redundant listeners

```
message LogEvent {  
    required uint64 pid = 1;  
    optional uint64 tid = 2;  
    optional uint64 sid = 4;  
    required uint64 sequence = 5;  
    required uint64 timestamp = 6;  
    enum Level { PANIC = 0, ERROR = 1.. }  
    required Level level = 7;  
    required bytes payload = 8;  
}
```

Announcements

- Dedicated channel.
- Payload: channel IP, channel port, last seq, pid, tid, sid + stats
- All announcers listen and self-throttle.
- Directory service accumulates

Consolidation

- Redundant journalers (RAID)
- ESP: detect loss in real time window
- If necessary, Map/Reduce processing to try to resolve partial loss.

Efficiency

- Hundreds of nodes
- >50MB/sec
- >50,000 pps
- 3-4 “journalers” resolving data
- >5TB reconciled data a day
- <0.1% data loss

Envelope 6

✿ Take out 6 envelopes . . .